

# Monitoring screenshots

Milestone 3 evidence: "Ops runbook (deploy/rollback/incident) in repo + **monitoring screenshots**". The three monitoring services — Sentry (errors), Grafana Cloud / Tempo (traces and metrics), PostHog (web analytics, no PII) — were instrumented in code (`apps/api/src/instrumentation.ts`, `apps/web/src/lib/observability.ts`) and verified live in production on 2026-09-08. The screenshots were captured on 2026-09-11, directly against the live dashboards.

Source: the working document `specs/EVIDENCIA-2026-09-11-monitoring-screenshots.md`, as of 2026-09-21.

**No staged data.** Every screenshot below shows what each service reported at capture time, unedited — real traffic and real errors from the production instance (`propnexus-api.onrender.com` / `propnexus-web.onrender.com`), not a mock or a seeded demo state.

## How this also covers Output 4's telemetry

Output 4 (Testing & Security) asks for "telemetry for coverage/latency/error budgets". The same instrumentation shown on this page provides two of the three, and the test report provides the third:


| Telemetry     | Where it comes from                                                                                      | Evidence                                                                    |
|---------------|----------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------|
| Coverage      | Vitest coverage on every CI run; the API has minimum thresholds that fail the build                      | The coverage table in the <a href="#">test report</a> (API: 89.3% of lines) |
| Latency       | OpenTelemetry traces of every API request, exported to Grafana Cloud (Tempo), with per-middleware timing | <code>grafana-tempo-trace-detail.jpg</code> below                           |
| Error budgets | Sentry captures every unhandled error in the API and the web app                                         | <code>sentry-issues.jpg</code> below                                        |

The reservation → escrow latency that acceptance criterion 3 measures (median under 12 minutes) is a separate, product-level metric, served by the API at `GET /api/v1/audit-logs/telemetry/reservation-to-escrow` and reported under evidence item 3.

**What is not in place yet:** no formal error-budget target (an SLO such as "99% of requests without a 5xx over 30 days") has been set. Sentry records the errors a budget would be computed from; setting the target itself is left for before mainnet.

## Files


All under [monitoring/](#).



Sentry's issue feed for the `propnexus-api` project, filtered to unresolved issues. Two real, unhandled errors are visible:


- **LibsqlError** – `SQLITE_CONSTRAINT: UNIQUE constraint failed: Unit.projectId, Unit.unitReference`, 1 hour old at capture time. This one is incidental, not a monitoring artifact: it was produced during this same session's manual exercise of the reservation-to-escrow flow, submitting a unit-creation form twice against the same reference. Left in the screenshot on purpose, as further proof the pipeline captures real errors as they happen, not synthetic ones.
- **Error** – `Cannot find module '@opentelemetry/api'`, 3 days old — a historical, already-fixed incident (see `apps/api/CLAUDE.md`, 2026-09-08 entry) whose issue entry is still open in Sentry because nobody has manually resolved it there; the underlying bug has been fixed in code since.

grafana-tempo-trace-detail.jpg



Grafana Cloud → Explore → the grafanacloud-giantmountain1601-traces (Tempo) datasource. Left panel: the trace search table over the last 24 hours, real propnex-api GET /health traces with real timestamps and durations. Right panel: one of those traces opened to its full waterfall — 18 spans, per-middleware timing breakdown (helmetMiddleware, jsonParser, etc.), 191.59 ms total duration, HTTP 200.


posthog-web-analytics.jpg



PostHog → Web analytics, last 7 days, for the apps/web production deployment. Shows a real visitor session against /login on 2026-09-10, with visitor/pageview/session counts, bounce rate, and the unique-

visitors time series. (The dedicated "Web vitals" tab returned no data for the selected range — traffic volume so far is too low for that specific metric to have recorded; the general Web analytics view already demonstrates the pipeline is live and receiving real events, which is what the evidence asks for.)

render-services-deployed.jpg



Not required on its own (the public URL is covered elsewhere), included as supporting context: both Render services, `propnexus-api` and `propnexus-web`, `Deployed` at capture time.

## Summary

| Service               | What it monitors                             | Status at capture                 |
|-----------------------|----------------------------------------------|-----------------------------------|
| Sentry                | Application errors ( <code>apps/api</code> ) | Live, 2 real unresolved issues    |
| Grafana Cloud (Tempo) | Distributed traces ( <code>apps/api</code> ) | Live, real traces in the last 24h |
| PostHog               | Web analytics ( <code>apps/web</code> )      | Live, real session recorded       |
| Render                | Deploy status (both services)                | <code>Deployed</code> , both      |